

1 Pre-Check

This section is designed as a conceptual check for you to determine if you conceptually understand and have any misconceptions about this topic. Please answer true/false to the following questions, and include an explanation:

- 1.1 True or False: C is a pass-by-value language.

True. If you want to pass a reference to anything, you should use a pointer.

- 1.2 What is a pointer? What does it have in common to an array variable?

As we like to say, "everything is just bits." A pointer is just a sequence of bits, interpreted as a memory address. An array acts like a pointer to the first element in the allocated memory for that array. However, an array name is not a variable, that is, `arr = arr` whereas `ptr != ptr` unless some magic happens (what does that mean?).

- 1.3 If you try to dereference a variable that is not a pointer, what will happen? What about when you free one?

It will treat that variable's underlying bits as if they were a pointer and attempt to access the data there. C will allow you to do almost anything you want, though if you attempt to access an "illegal" memory address, it will segfault for reasons we will learn later in the course. It's why C is not considered "memory safe": you can shoot yourself in the foot if you're not careful. If you free a variable that either has been freed before or was not malloced/calloced/reallocated, bad things happen. The behavior is undefined and terminates execution, resulting in an "invalid free" error.

- 1.4 When should you use the heap over the stack? Do they grow?

If you need to keep access to data over several function calls, use the heap. If you're dealing with a large piece of data, passing around a pointer to something on the heap is more efficient and a better practice than passing around the data itself. (Think: carrying a library around vs knowing the address TO the library). Heaps grow up and stacks grow down, meeting when working memory is full.

2 C


C is syntactically similar to Java, but there are a few key differences:

1. C is function-oriented, not object-oriented; there are no objects.
2. C does not automatically handle memory for you.

- Stack memory, or *things that are not manually allocated*: data is garbage immediately after the *function in which it was defined* returns.
 - Heap memory, or *things allocated with malloc, calloc, or realloc*: data is freed only when the programmer explicitly frees it!
 - There are two other sections of memory that we learn about in this course, *static* and *code*, but we'll get to those later.
 - In any case, allocated memory always holds garbage until it is initialized!
3. C uses pointers explicitly. If `p` is a pointer, then `*p` tells us to use the value that `p` points to, rather than the value of `p`, and `&x` gives the address of `x` rather than the value of `x`.

On the left is the memory represented as a box-and-pointer diagram.

On the right, we see how the memory is really represented in the computer.



Let's assume that `int* p` is located at `0xF9320904` and `int x` is located at `0xF93209B0`. As we can observe:

- `*p` evaluates to `0x2A` (42_{10}).
- `p` evaluates to `0xF93209AC`.
- `x` evaluates to `0x61C`.
- `&x` evaluates to `0xF93209B0`.

Let's say we have an `int **pp` that is located at `0xF9320900`.

2.1 What does `pp` evaluate to? How about `*pp`? What about `**pp`?

`pp` evaluates to `0xF9320904`. `*pp` evaluates to `0xF93209AC`. `**pp` evaluates to `0x2A`.

2.2 The following functions are syntactically-correct C, but written in an incomprehensible style. Describe the behavior of each function in plain English.

- (a) Recall that the ternary operator evaluates the condition before the `?` and returns the value before the colon (`:`) if true, or the value after it if false.

```

1  int foo(int *arr, size_t n) {
2      return n ? arr[0] + foo(arr + 1, n - 1) : 0;
3  }
```

Returns the sum of the first N elements in `arr`.

- (b) Recall that the negation operator, `!`, returns 0 if the value is non-zero, and 1 if the value is 0. The `~` operator performs a *bitwise not* (NOT) operation.

```

1  int bar(int *arr, size_t n) {
2      int sum = 0, i;
3      for (i = n; i > 0; i--)
4          sum += !arr[i - 1];
5      return ~sum + 1;
6  }
```



Returns -1 times the number of zeroes in the first N elements of `arr`.

- (c) Recall that `^` is the *bitwise exclusive-or* (XOR) operator.

```

1  void baz(int x, int y) {
2      x = x ^ y;
3      y = x ^ y;
4      x = x ^ y;
5  }
```

Ultimately does not change the value of either `x` or `y`.

- (d) (Bonus: How do you write the *bitwise exclusive-nor* (XNOR) operator in C?)

```
1  x == y
```

3 Programming with Pointers

3.1 Implement the following functions so that they work as described.

- (a) Swap the value of two `ints`. *Remain swapped after returning from this function.*

```

1  void swap(int *x, int *y) {
2      int temp = *x;
3      *x = *y;
4      *y = temp;
5  }
```

- (b) Return the number of bytes in a string. *Do not use `strlen`.*

```

1  int mystrlen(char* str) {
2      int count = 0;
3      while (*str++) {
4          count++;
5      }
```

```

6     return count;
7 }

```

3.2 The following functions may contain logic or syntax errors. Find and correct them.

- (a) Returns the sum of all the elements in `summands`.

It is necessary to pass a size alongside the pointer.

```

1 int sum(int* summands, size_t n) {
2     int sum = 0;
3     for (int i = 0; i < n; i++)
4         sum += *(summands + i);
5     return sum;
6 }

```

- (b) Increments all of the letters in the `string` which is stored at the front of an array of arbitrary length, `n >= strlen(string)`. Does not modify any other parts of the array's memory.

The ends of strings are denoted by the null terminator rather than *n*. Simply having space for *n* characters in the array does not mean the string stored inside is also of length *n*.

```

1 void increment(char* string) {
2     for (i = 0; string[i] != 0; i++)
3         string[i]++; // or (*(string + i))++;
4 }

```

Another common bug to watch out for is the corner case that occurs when incrementing the character with the value `0xFF`. Adding 1 to `0xFF` will overflow back to 0, producing a null terminator and unintentionally shortening the string.

- (c) Copies the string `src` to `dst`.

```

1 void copy(char* src, char* dst) {
2     while (*dst++ = *src++);
3 }

```

No errors.

- (d) Overwrites an input string `src` with "61C is awesome!" if there's room. Does nothing if there is not. Assume that `length` correctly represents the length of `src`.

```

1 void cs61c(char* src, size_t length) {
2     char *srcptr, replaceptr;

```

```

3     char replacement[16] = "61C is awesome!";
4     srcptr = src;
5     replaceptr = replacement;
6     if (length >= 16) {
7         for (int i = 0; i < 16; i++)
8             *srcptr++ = *replaceptr++;
9     }
10 }

```

`char *srcptr`, `replaceptr` initializes a `char` pointer, and a `char`—not two `char` pointers.

The correct initialization should be, `char *srcptr`, `*replaceptr`.

4 Memory Management

4.1 For each part, choose one or more of the following memory segments where the data could be located: **code**, **static**, **heap**, **stack**.

(a) Static variables

Static

(b) Local variables

Stack

(c) Global variables

Static

(d) Constants

Code, static, or stack

Constants can be compiled directly into the code. `x = x + 1` can compile with the number 1 stored directly in the machine instruction in the code. That instruction will always increment the value of the variable `x` by 1, so it can be stored directly in the machine instruction without reference to other memory. This can also occur with pre-processor macros.

```

1  #define y 5
2
3  int plus_y(int x) {
4      x = x + y;
5      return x;
6  }

```

Constants can also be found in the stack or static storage depending on if it's declared in a function or not.

```

1  const int x = 1;
2
3  int sum(int* arr) {
4      int total = 0;
5      ...
6  }
```

In this example, `x` is a variable whose value will be stored in the static storage, while `total` is a local variable whose value will be stored on the stack. Variables declared **const** are not allowed to change, but the usage of **const** can get more tricky when combined with pointers.

(e) Machine Instructions

Code

(f) Result of `malloc`

Heap

(g) String Literals

Static.

When declared in a function, string literals can only be stored in static (or data) memory. String literals are declared when a character pointer is assigned to a string declared within quotation marks, i.e. `char* s = "string"`. Strings declared this way are stored in the static memory segment. You'll often see 3 other ways in which to declare strings, one being very similar to string literals, as following: `char[7] s = "string"`. This is a string array will be stored in the stack and is mutable. Note that the compiler will arrange for the char array to be initialised from the literal and be mutable.

4.2 Write the code necessary to allocate memory on the heap in the following scenarios

(a) An array `arr` of k integers

```
arr = (int *) malloc(sizeof(int) * k);
```

(b) A string `str` containing p characters

```
str = (char *) malloc(sizeof(char) * (p + 1));
```

Don't forget the null terminator!

(c) An $n \times m$ matrix `mat` of integers initialized to zero.

```
mat = (int *) calloc(n * m, sizeof(int));
```

Alternative solution. This might be needed if you wanted to efficiently permute the rows of the matrix.

```

1  mat = (int **) calloc(n, sizeof(int *));
2  for (int i = 0; i < n; i++)
3      mat[i] = (int *) calloc(m, sizeof(int));
```

- 4.3 What's the main issue with the code snippet seen here? (Hint: `gets()` is a function that reads in user input and stores it in the array given in the argument.)

```

1 char* foo() {
2     char buffer[64];
3     gets(buffer);
4
5     char* important_stuff = (char*) malloc(11 * sizeof(char));
6
7     int i;
8     for (i = 0; i < 10; i++) important_stuff[i] = buffer[i];
9     important_stuff[i] = '\0';
10    return important_stuff;
11 }

```

If the user input contains more than 63 characters, then the input will override other parts of the memory! (You will learn more about this and how it can be used to maliciously exploit programs in [CS 161](#).)

Note that it's perfectly acceptable in C to create an array on the stack. It's often discouraged (mostly because people often forget the array was initialized on the stack and accidentally return a pointer to it), but it's not an issue itself.

Suppose we've defined a linked list **struct** as follows. Assume `*lst` points to the first element of the list, or is `NULL` if the list is empty.

```

struct ll_node {
    int first;
    struct ll_node* rest;
}

```

- 4.4 Implement `prepend`, which adds one new value to the front of the linked list. Hint: why use `ll_node **lst` instead of `ll_node*lst`?

```

1 void prepend(struct ll_node** lst, int value) {
2     struct ll_node* item = (struct ll_node*) malloc(sizeof(struct ll_node));
3     item->first = value;
4     item->rest = *lst;
5     *lst = item;
6 }

```

- 4.5 Implement `free_ll`, which frees all the memory consumed by the linked list.

```

1 void free_ll(struct ll_node** lst) {
2     if (*lst) {

```

```
3     free_ll(&((*lst)->rest));
4     free(*lst);
5 }
6 *lst = NULL; // Make writes to **lst fail instead of writing to unusable memory.
7 }
```