

UC Berkeley  
Teaching Professor  
Dan Garcia

# CS61C

## Great Ideas in Computer Architecture (a.k.a. Machine Structures)



UC Berkeley  
Professor  
Bora Nikolić

### Caches I

# Binary Prefix



# Kilo, Mega, Giga, Tera, Peta, Exa, Zetta, Yotta

- Common use prefixes (all SI, except K [= k in SI])
- Confusing! Common usage of "kilobyte" means 1024 bytes, but the "correct" SI value is 1000 bytes
- Hard Disk manufacturers & Telecommunications are the only computing groups that use SI factors
  - What is advertised as a 1 TB drive actually holds about 90% of what you expect
  - A 1 Mbit/s connection transfers  $10^6$  bps.

Name	Abbr	Factor	SI size
Kilo	K	$2^{10} = 1,024$	$10^3 = 1,000$
Mega	M	$2^{20} = 1,048,576$	$10^6 = 1,000,000$
Giga	G	$2^{30} = 1,073,741,824$	$10^9 = 1,000,000,000$
Tera	T	$2^{40} = 1,099,511,627,776$	$10^{12} = 1,000,000,000,000$
Peta	P	$2^{50} = 1,125,899,906,842,624$	$10^{15} = 1,000,000,000,000,000$
Exa	E	$2^{60} = 1,152,921,504,606,846,976$	$10^{18} = 1,000,000,000,000,000,000$
Zetta	Z	$2^{70} = 1,180,591,620,717,411,303,424$	$10^{21} = 1,000,000,000,000,000,000,000$
Yotta	Y	$2^{80} = 1,208,925,819,614,629,174,706,176$	$10^{24} = 1,000,000,000,000,000,000,000,000$



# kibi, mebi, gibi, tebi, pebi, exbi, zebi, yobi

- IEC Standard Prefixes [only to exbi officially]

Name	Abbr	Factor
kibi	Ki	$2^{10} = 1,024$
mebi	Mi	$2^{20} = 1,048,576$
gibi	Gi	$2^{30} = 1,073,741,824$
tebi	Ti	$2^{40} = 1,099,511,627,776$
pebi	Pi	$2^{50} = 1,125,899,906,842,624$
exbi	Ei	$2^{60} = 1,152,921,504,606,846,976$
zebi	Zi	$2^{70} = 1,180,591,620,717,411,303,424$
yobi	Yi	$2^{80} = 1,208,925,819,614,629,174,706,176$

- International Electrotechnical Commission (IEC) in 1999 introduced these to specify binary quantities.
- Names come from shortened versions of the original SI prefixes (same pronunciation) and *bi* is short for "binary", but pronounced "bee" :-)
- Now SI prefixes only have their base-10 meaning and never have a base-2 meaning.


# Kilo, Mega, Giga, Tera, Peta, Exa, Zetta, Yotta

1. Kid meets giant Texas people exercising zen-like yoga. – Rolf O
2. Kind men give ten percent extra, zestfully, youthfully. – Hava E
3. Kissing Mentors Gives Testy Persistent Extremists Zealous Youthfulness. – Gary M
4. Kindness means giving, teaching, permeating excess zeal yourself. – Hava E
5. Killing messengers gives terrible people exactly zero, yo
6. Kindergarten means giving teachers perfect examples (of) zeal (&) youth
7. Kissing mediocre giraffes teaches people (to) expect zero (from) you
8. Kinky Mean Girls Teach People Exciting Zen Yoga
9. Kissing Mel Gibson, Teddy Pendergrass exclaimed: "Zesty, yo!" – Dan G
10. Kissing me gives ten percent extra zeal & youth! – Dan G (borrowing parts)

# The way to remember #s

- What is  $2^{34}$ ? How many bits to address (i.e., what's  $\lceil \log_2 = \lg \text{of} \rceil$ ) 2.5 TiB?
- Answer!  $2^{XY}$  means...

X=0 $\Rightarrow$ ---	Y=0 $\Rightarrow$ 1
X=1 $\Rightarrow$ kibi $\sim 10^3$	Y=1 $\Rightarrow$ 2
X=2 $\Rightarrow$ mebi $\sim 10^6$	Y=2 $\Rightarrow$ 4
X=3 $\Rightarrow$ gibi $\sim 10^9$	Y=3 $\Rightarrow$ 8
X=4 $\Rightarrow$ tebi $\sim 10^{12}$	Y=4 $\Rightarrow$ 16
X=5 $\Rightarrow$ pebi $\sim 10^{15}$	Y=5 $\Rightarrow$ 32
X=6 $\Rightarrow$ exbi $\sim 10^{18}$	Y=6 $\Rightarrow$ 64
X=7 $\Rightarrow$ zebi $\sim 10^{21}$	Y=7 $\Rightarrow$ 128
X=8 $\Rightarrow$ yobi $\sim 10^{24}$	Y=8 $\Rightarrow$ 256
	Y=9 $\Rightarrow$ 512



MEMORIZE!

# Library Analogy

# New-School Machine Structures

Software

Parallel Requests

Assigned to computer

e.g., Search "Cats"

Parallel Threads

Assigned to core e.g., Lookup, Ads

Parallel Instructions

>1 instruction @ one time

e.g., 5 pipelined instructions

Parallel Data

>1 data item @ one time

e.g., Add of 4 pairs of words

Hardware descriptions

All gates work in parallel at same time


Harness Parallelism & Achieve High Performance

Hardware


Warehouse Scale Computer



Smart Phone



# Components of a Computer




# Why are Large Memories Slow? Library Analogy

- Time to find a book in a large library
  - Search a large card catalog – (mapping title/author to index number)
  - Round-trip time to walk to the stacks and retrieve the desired book
- Larger libraries worsen both delays
- Electronic memories have same issue, *plus* the technologies used to store a bit slow down as density increases (e.g., SRAM vs. DRAM vs. Disk)



What we want is a large, yet fast memory!

# Processor-DRAM Gap (Latency)



1980 microprocessor executes ~one instruction in same time as DRAM access  
2020 microprocessor executes ~1000 instructions in same time as DRAM access

**Slow DRAM access has disastrous impact on CPU performance!**

# Memory Hierarchy

# What To Do: Library Analogy


- Write a report using library books
  - E.g., works of J.D. Salinger
- Go to library, look up books, fetch from stacks, and place on desk in library
- If need more, check out, keep on desk
  - But don't return earlier books since might need them
- You hope this collection of ~10 books on desk enough to write report, despite 10 being only 0.00001% of books in UC Berkeley libraries




# Memory Caching

- Mismatch between processor and memory speeds leads us to add a new level...
  - Introducing a “memory cache”
- Implemented with same IC processing technology as the CPU (usually integrated on same chip)
  - faster but more expensive than DRAM memory.
- Cache is a copy of a subset of main memory
- Most processors have separate caches for instructions and data.


# Great Idea #3: Principle of Locality / Memory Hierarchy



# Jim Gray's Storage Latency Analogy: How Far Away is the Data?




# Characteristics of the Memory Hierarchy



# Typical Memory Hierarchy

- **The Trick:** present processor with as much memory as is available in the *cheapest* technology at the speed offered by the *fastest* technology



# Memory Hierarchy

- If level closer to Processor, it is:
  - Smaller
  - Faster
  - More expensive
  - subset of lower levels (contains most recently used data)
- Lowest Level (usually disk=HDD/SSD) contains all available data (does it go beyond the disk?)
- Memory Hierarchy presents the processor with the **illusion** of a very large & fast memory



Locality, Design,  
Management

# Memory Hierarchy Basis

- Cache contains copies of data in memory that are being used.
- Memory contains copies of data on disk that are being used.
- Caches work on the principles of **temporal** and **spatial locality**.
  - Temporal locality (locality in time): If we use it now, chances are we'll want to use it again soon.
  - Spatial locality (locality in space): If we use a piece of memory, chances are we'll use the neighboring pieces soon.

# What to Do About Locality

- *Temporal Locality*
  - If a memory location is referenced then it will tend to be referenced again soon
    - ⇒ Keep most recently accessed data items closer to the processor
- *Spatial Locality*
  - If a memory location is referenced, the locations with nearby addresses will tend to be referenced soon
    - ⇒ Move blocks consisting of contiguous words closer to the processor

# Cache Design

- How do we organize cache?
- Where does each memory address map to?
  - (Remember that cache is subset of memory, so multiple memory addresses map to the same cache location.)
- How do we know which elements are in cache?
- How do we quickly locate them?

# How is the Hierarchy Managed?

- registers  $\leftrightarrow$  memory
  - By compiler (or assembly level programmer)
- cache  $\leftrightarrow$  main memory
  - By the cache controller hardware
- main memory  $\leftrightarrow$  disks (secondary storage)
  - By the operating system (virtual memory)
  - Virtual to physical address mapping assisted by the hardware ('translation lookaside buffer' or TLB)
  - By the programmer (files)      **Also a type of cache**



# “And in Conclusion...”

- Caches provide an illusion to the processor that the memory is infinitely large and infinitely fast

